

Linux Strace Command



Name	:	Sajjad Ali
CMSID	:	023-19-0100
Department	:	Computer Science
Semester	:	BS IV Section B
Course	:	Operating System (Theory)
Assignment	:	Strace (#01)
Submitted to	:	Dr. Raheel Memon

S trace

It is a diagnostic, debugging and instructional user space utility for Linux. It is used to monitor and temper with interactions processes and Linux kernel, which include system calls, signal deliveries and changes of process state.

It can be installed through this command on terminal:

```
$ sudo apt install strace
```

- We can check the commands in strace via

```
$ strace -h
```

- We can trace any running processor via following command:

```
$ sudo strace -c -p PID
```

where PID is the processor Id of any running processor.

Suppose we have run following command:

```
$ sudo strace -c -p 2010
```

The most probable output is the following type:

```

genius@192:~$ sudo strace -c -p 2010
strace: Process 2010 attached
^Cstrace: Process 2010 detached

```

% time	seconds	uscs/call	calls	errors	syscall
34.43	0.001081	5	207	176	recvmsg
24.17	0.000759	9	83		poll
17.32	0.000544	9	59		ioctl
10.67	0.000335	8	39		writew
6.56	0.000206	5	40		getpid
4.90	0.000154	5	29		write
1.59	0.000050	4	12		read
0.35	0.000011	11	1		restart_syscall
100.00	0.003140		470	176	total

Here syscall refers to the System Calls.

- We can also trace with respect to some predicate or condition.

\$ sudo strace -q -e trace=process df -h

```

genius@192:~$ sudo strace -q -e trace=process df -h
execve("/usr/bin/df", ["df", "-h"], 0x7ffc80957ba0 /* 25 vars */) = 0
arch_prctl(0x3001 /* ARCH_??? */, 0x7fff97c69860) = -1 EINVAL (Invalid argument)
arch_prctl(ARCH_SET_FS, 0x7f4dc21f9580) = 0
Filesystem      Size  Used Avail Use% Mounted on
udev            1.9G   0    1.9G   0% /dev
tmpfs           385M  2.0M  383M   1% /run
/dev/sda7       56G   25G   28G  48% /
tmpfs           1.9G  4.0K   1.9G   1% /dev/shm
tmpfs           5.0M  4.0K   5.0M   1% /run/lock
tmpfs           1.9G   0    1.9G   0% /sys/fs/cgroup
/dev/loop1      128K  128K   0 100% /snap/anbox-installer/24
/dev/loop3      56M   56M   0 100% /snap/core18/1988
/dev/loop4     256K  256K   0 100% /snap/figlet/36
/dev/loop0      99M   99M   0 100% /snap/core/10823
/dev/loop5     163M  163M   0 100% /snap/gnome-3-28-1804/145
/dev/loop7      65M   65M   0 100% /snap/gtk-common-themes/1514
/dev/loop6     219M  219M   0 100% /snap/gnome-3-34-1804/66
/dev/loop10     32M   32M   0 100% /snap/snapd/11036
/dev/loop8      52M   52M   0 100% /snap/snap-store/518
/dev/loop11     33M   33M   0 100% /snap/snapd/11107
/dev/loop9     425M  425M   0 100% /snap/pycharm-community/226
/dev/loop2     100M  100M   0 100% /snap/core/10859
/dev/sda1       96M   33M   64M  35% /boot/efi
tmpfs           385M  100K  384M   1% /run/user/1000
exit_group(0)
+++ exited with 0 +++

```

Like this we can also trace for file, memory, network or signal just by replacing process with others.

- We can also save the result of trace in any file.

\$ sudo strace -o filename df -h

```
gentus@192:~/Downloads$ sudo strace -o file.txt df -h
Filesystem            Size  Used Avail Use% Mounted on
udev                  1.9G     0  1.9G   0% /dev
tmpfs                  385M   2.0M  383M   1% /run
/dev/sda7              56G   25G   28G  48% /
tmpfs                  1.9G   4.0K   1.9G   1% /dev/shm
tmpfs                  5.0M   4.0K   5.0M   1% /run/lock
tmpfs                  1.9G     0  1.9G   0% /sys/fs/cgroup
/dev/loop1            128K  128K     0 100% /snap/anbox-installer/24
/dev/loop3             56M   56M     0 100% /snap/core18/1988
/dev/loop4            256K  256K     0 100% /snap/figlet/36
/dev/loop0             99M   99M     0 100% /snap/core/10823
/dev/loop5            163M  163M     0 100% /snap/gnome-3-28-1804/145
/dev/loop7             65M   65M     0 100% /snap/gtk-common-themes/1514
/dev/loop6            219M  219M     0 100% /snap/gnome-3-34-1804/66
/dev/loop10            32M   32M     0 100% /snap/snapd/11036
/dev/loop8             52M   52M     0 100% /snap/snap-store/518
/dev/loop11            33M   33M     0 100% /snap/snapd/11107
/dev/loop9            425M  425M     0 100% /snap/pycharm-community/226
/dev/loop2            100M  100M     0 100% /snap/core/10859
/dev/sda1              96M   33M   64M  35% /boot/efi
tmpfs                  385M  100K  384M   1% /run/user/1000
```

- We can also trace the programs written in any language whether we have its source code or not.

Suppose hello.c program is written in c and prints hello world on the console. After compiling, we are tracing this program via:

\$ strace -c ./hello

Hello World%	time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	1			read
0.00	0.000000	0	1			write
0.00	0.000000	0	2			close
0.00	0.000000	0	3			fstat
0.00	0.000000	0	7			mmap
0.00	0.000000	0	4			mprotect
0.00	0.000000	0	1			munmap
0.00	0.000000	0	3			brk
0.00	0.000000	0	6			pread64
0.00	0.000000	0	1	1		access
0.00	0.000000	0	1			execve
0.00	0.000000	0	2	1		arch_prctl
0.00	0.000000	0	2			openat
100.00	0.000000		34	2		total

Similarly tracing the java program named hi.java after compiling,

\$ strace -c java hi

h

Hi	% time	seconds	usecs/call	calls	errors	syscall
	95.47	0.006469	3234	2		futex
	1.48	0.000100	2	45		mmap
	0.81	0.000055	3	16		mprotect
	0.63	0.000043	43	1		clone
	0.56	0.000038	12	3		munmap
	0.32	0.000022	0	41	29	openat
	0.21	0.000014	1	11		read
	0.18	0.000012	0	13		close
	0.18	0.000012	1	12		fstat
	0.16	0.000011	5	2		getpid
	0.00	0.000000	0	24	21	stat
	0.00	0.000000	0	3		brk
	0.00	0.000000	0	2		rt_sigaction
	0.00	0.000000	0	1		rt_sigprocmask
	0.00	0.000000	0	8		pread64
	0.00	0.000000	0	2	1	access
	0.00	0.000000	0	1		execve
	0.00	0.000000	0	2		readlink
	0.00	0.000000	0	2	1	arch_prctl
	0.00	0.000000	0	1		set_tid_address
	0.00	0.000000	0	1		set_robust_list
	0.00	0.000000	0	1		prlimit64
	100.00	0.006776		194	52	total

We can see that java program makes more system calls than the same C based program.